

PROJECT : HEALTH INSURANCE PREDICTION

- DATE : 19.03.2024
- NAME : yashwanth
- DOMAIN : MACHINE LEARNIG
 - DEGREE : BTECH AI AND DS (3RD YEAR)

✓ DATA SET

```
#importing libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as py
import seaborn as sb
```

```
#reading the dataset (DS)
```

```
med_ins=pd.read_csv("insurance.csv")
print(med_ins)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

```
[1338 rows x 7 columns]
```

```
print(med_ins.nunique())
```

```
age          47
sex           2
bmi          548
children      6
smoker        2
region        4
charges     1337
dtype: int64
```

```
print(med_ins["region"].unique())
```

```
['southwest' 'southeast' 'northwest' 'northeast']
```

```
#information of DS
```

```
print(med_ins.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1338 non-null   int64   
 1   sex         1338 non-null   object  
 2   bmi         1338 non-null   float64  
 3   children    1338 non-null   int64   
 4   smoker      1338 non-null   object  
 5   region      1338 non-null   object  
 6   charges     1338 non-null   float64  
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
None
```

```
#checking null values in DS
```

```
print(med_ins.isna().sum())
```

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

```
#displaying the columns
```

```
print(med_ins.columns)
```

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

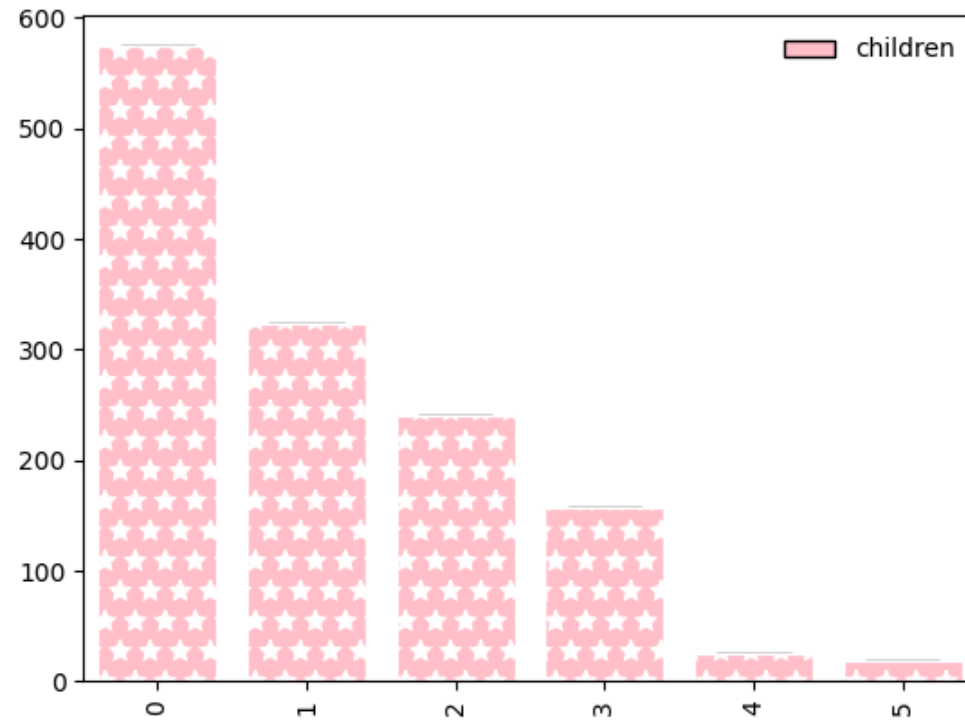
```
py.figure(figsize=(5,8))
med_ins['age'].value_counts().plot(kind='pie',colormap='RdPu')
py.legend(title='Ages',bbox_to_anchor=(1.05,0.75),ncol=10,loc='best',frameon=False,prop={'weight':'bold'},fontsize=6)
```

```
<matplotlib.legend.Legend at 0x7ddf06236860>
```




```
a=med_ins['children'].value_counts()
a.plot(kind='bar',color='pink',edgecolor='black')
py.legend(loc='best',frameon=False)
py.bar(range(len(a)),a,hatch='*',color='pink',edgecolor='white')
```

<BarContainer object of 6 artists>

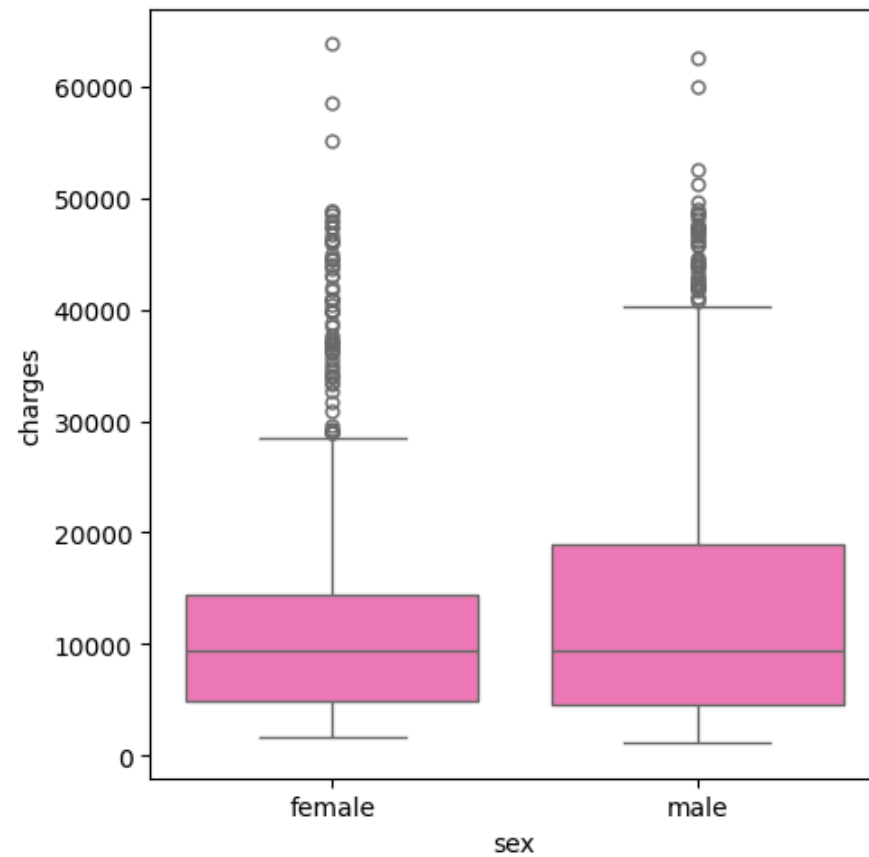


```
b=med_ins['bmi'].value_counts().plot(kind='hist',color='pink',edgecolor='black')
ax1=py.gca()
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
py.legend(loc='best',frameon=False,bbox_to_anchor=(0.9,0.99))
```

<matplotlib.legend.Legend at 0x7ddf06235ff0>




```
sb.catplot(x="sex",y="charges",kind='box',color='hotpink',data=med_ins)
#ax=py.axes()
#ax.set_facecolor("black")
ax1=py.gca()
ax1.spines['top'].set_visible(True)
ax1.spines['right'].set_visible(True)
```



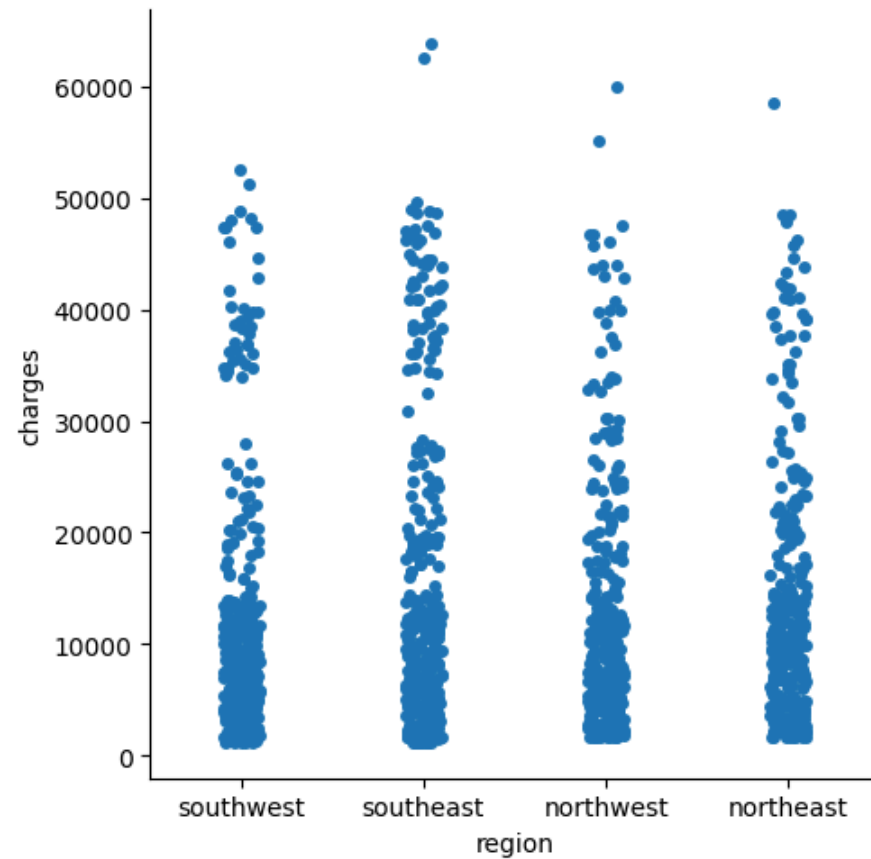
```
sb.violinplot(x="sex",y="bmi",hue="smoker",data=med_ins)  
py.legend(bbox_to_anchor=(1.17,0.17),frameon=True)
```

<matplotlib.legend.Legend at 0x7ddf03fdf010>



```
sb.catplot(x="region",y="charges",data=med_ins)
```

<seaborn.axisgrid.FacetGrid at 0x7ddf0236a170>



#decription of DS

```
print(med_ins.describe())
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010


```

from sklearn.preprocessing import LabelEncoder
l=LabelEncoder()
med_ins.iloc[:,1]=l.fit_transform(med_ins.iloc[:,1])

med_ins.iloc[:,4]=l.fit_transform(med_ins.iloc[:,4])

med_ins.iloc[:,5]=l.fit_transform(med_ins.iloc[:,5])

print(med_ins)

```

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520
...
1333	50	1	30.970	3	0	1	10600.54830
1334	18	0	31.920	0	0	0	2205.98080
1335	18	0	36.850	0	0	2	1629.83350
1336	21	0	25.800	0	0	3	2007.94500
1337	61	0	29.070	0	1	1	29141.36030

[1338 rows x 7 columns]

```

<ipython-input-15-3b6b98c751e9>:3: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values inplace
med_ins.iloc[:,1]=l.fit_transform(med_ins.iloc[:,1])
<ipython-input-15-3b6b98c751e9>:5: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values inplace
med_ins.iloc[:,4]=l.fit_transform(med_ins.iloc[:,4])
<ipython-input-15-3b6b98c751e9>:7: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values inplace
med_ins.iloc[:,5]=l.fit_transform(med_ins.iloc[:,5])

```

✓ LinearRegression

```

#assign x and y

x=med_ins[['smoker']].values
y=med_ins[['charges']].values

```

```
#spliting
```

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
#algorithm
```


```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr1=lr.fit(x_train,y_train)  
print('COEFFICIENT:',lr1.coef_)  
print(' INTERCEPT:',lr1.intercept_)
```

```
COEFFICIENT: [[23332.42129457]]  
INTERCEPT: [8354.09098751]
```

```
#prediction
```

```
y_pred=lr1.predict(x_test)  
ax=py.axes()  
ax1=py.gca()  
ax.set_facecolor("white")  
ax1.spines['top'].set_visible(False)  
ax1.spines['right'].set_visible(False)  
py.scatter(x_test,y_test,s=30,color="pink")  
py.plot(x_test,y_pred,color="black",linewidth=4,linestyle=(0,(0.1,2)),dash_capstyle="round")
```


[<matplotlib.lines.Line2D at 0x7ddefefd1f30>]



#Difference prediction

```
res=y_test-y_pred
y=[i for i in range(1,len(res)+1)]
ax=py.axes()
ax1=py.gca()
ax.set_facecolor("white")
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
py.scatter(y,res,s=30,color='pink')
py.plot(y,[0]*len(y_pred),color='black',linewidth=4,linestyle=(0,(0.1,2)),dash_capstyle="round")
```

[<matplotlib.lines.Line2D at 0x7ddef047d60>]



```
from sklearn.metrics import r2_score, mean_squared_error
mse = mean_squared_error(y_test, y_pred)
print("ROOT MEAN SQUARED ERROR : ", np.sqrt(mse))
print("    MEAN SQUARED ERROR : ", mse)
A1 = r2_score(y_test, y_pred)
print("    ACCUACY SCORE : ", A1)
```

```
ROOT MEAN SQUARED ERROR : 7165.756777678182
    MEAN SQUARED ERROR : 51348070.19684079
    ACCUACY SCORE : 0.6406587599098937
```

✓ PRE PROCESSING

```
#assign x and y
```

```
x=med_ins.iloc[:,0:-1].values  
y=med_ins.iloc[:, -1].values
```

```
#split
```

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
#Normalize
```

```
from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
x_train=sc.fit_transform(x_train)  
x_test=sc.fit_transform(x_test)
```

✓ DECISION TREE REGRESSOR

```
#Algorithm
```

```
from sklearn.tree import DecisionTreeRegressor  
dtr = DecisionTreeRegressor()  
dtr1 = dtr.fit(x_train,y_train)
```

```
#Predict
```

```
y_pred = dtr1.predict(x_test)  
from sklearn.metrics import r2_score,mean_squared_error  
a1=mean_squared_error(y_test,y_pred)  
r1=print('ROOT MEAN SQUARED ERROR : ',np.sqrt(a1))  
print("      MEAN SQUARED ERROR : ",a1)  
B1_=r2_score(y_test,y_pred)  
print('      ACCUACY SCORE : ',B1_)
```

```
ROOT MEAN SQUARED ERROR : 6680.25433580387
MEAN SQUARED ERROR : 44625797.99102641
ACCURACY SCORE : 0.6811991216237633
```

✓ NAIVE BAYES

```
#Algorithm
```

```
from sklearn.linear_model import BayesianRidge
BR=BayesianRidge()
BR2=BR.fit(x_train,y_train)
```

```
#predict
```

```
y_pred=BR2.predict(x_test)
from sklearn.metrics import r2_score,mean_squared_error
a1=mean_squared_error(y_test,y_pred)
r1=print('ROOT MEAN SQUARED ERROR : ',np.sqrt(a1))
print("      MEAN SQUARED ERROR : ",a1)
C1=r2_score(y_test,y_pred)
print('      ACCURACY SCORE : ',C1_)
```

```
ROOT MEAN SQUARED ERROR : 5976.886019951411
MEAN SQUARED ERROR : 35723166.49549062
ACCURACY SCORE : 0.7447983594728539
```

✓ RANDOM FOREST TREE

```
#ALGORITHM
```

```
from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor()
rf1=rf.fit(x_train,y_train)
```

```
#predict

y_pred=rf1.predict(x_test)
from sklearn.metrics import r2_score,mean_squared_error
a1=mean_squared_error(y_test,y_pred)
r1=print('ROOT MEAN SQUARED ERROR : ',np.sqrt(a1))
print("      MEAN SQUARED ERROR : ",a1)
D1=r2_score(y_test,y_pred)
print('      ACCUACY SCORE : ',D1_)
```

```
ROOT MEAN SQUARED ERROR :  4707.5588644641
MEAN SQUARED ERROR :  22161110.462394528
ACCUACY SCORE : 0.8416839182881424
```

✓ K-NEAREST NEIGHBOUR

```
#ALGORITHM
```

```
from sklearn.neighbors import KNeighborsRegressor
kn=KNeighborsRegressor(n_neighbors=5,metric='minkowski',p=2)
kn1=kn.fit(x_train,y_train)
```

```
#predict

y_pred=kn1.predict(x_test)
from sklearn.metrics import r2_score,mean_squared_error
a1=mean_squared_error(y_test,y_pred)
r1=print('ROOT MEAN SQUARED ERROR : ',np.sqrt(a1))
print("      MEAN SQUARED ERROR : ",a1)
E1=r2_score(y_test,y_pred)
print('      ACCUACY SCORE : ',E1_)
```

```

ROOT MEAN SQUARED ERROR : 5092.860629799064
MEAN SQUARED ERROR : 25937229.394557316
ACCURACY SCORE : 0.8147078173191759

```

✓ SUPPORT VECTOR MACHINE (SVM)

```
#ALGORITHM
```

```

from sklearn.svm import SVR
sv=SVR()
sv1=sv.fit(x_train,y_train)

```

```

#predict
y_pred=sv1.predict(x_test)

from sklearn.metrics import r2_score,mean_squared_error
a1=mean_squared_error(y_test,y_pred)
r1=print('ROOT MEAN SQUARED ERROR : ',np.sqrt(a1))
print("      MEAN SQUARED ERROR : ",a1)
F1_=r2_score(y_test,y_pred)
print('      ACCURACY SCORE : ',F1_)

```

```

ROOT MEAN SQUARED ERROR : 12719.78303542988
MEAN SQUARED ERROR : 161792880.46840978
ACCURACY SCORE : -0.15582722842788033

```

✓ FEATURE IMPORTANCE

```

lst=[A1_,B1_,C1_,D1_,E1_,F1_]
print(          )
print('          ALGORITHM ACCURACY SCORES          ')
print(          )

```



```

print(      )
print("LINEAR REGRESSION ACCUARCY SCORE : ",A1_)
print("RANDOM FOREST ACCURACY SCORE      : ",D1_)
print("DECISION TREE ACCURACY SCORE      : ",B1_)
print("NAIVE BAYES ACCURACY SCORE         : ",C1_)
print("KNN ACCURACY SCORE                 : ",E1_)
print("SVM ACCURACY SCORE                  : ",F1_)
print(      )
print(      )
print(      )

max_=max(lst)

if max_==A1_:

    print('SELECTED ALGORITHM : LINEAR ')
    print('TYPE OF ALGORITHM   : REGRESSION')
    print('ACCURACY SCORE      : ',max_)
    print(      )
    print('CONTRIBUTION')
    print(      )

    importance=lr.feature_importances_
    for i,v in enumerate(importance):
        print("Feature: %0d ; Score: %5f"%(i,v))

    index=med_ins.columns[:-1]
    importance=pd.Series(rf.feature_importances_,index=index)
    importance.nlargest(13).plot(kind='barh',color='pink')

elif max_==B1_:

    print('SELECTED ALGORITHM : DECISION TREE: ')
    print('TYPE OF ALGORITHM   : REGRESSION')
    print('ACCURACY SCORE      : ',max_)
    print(      )
    print('CONTRIBUTION')
    print(      )

    importance1=dtr.feature_importances_
    for i,v in enumerate(importance1):
        print("Feature: %0d ; Score: %5f"%(i,v))

```

```
index1=med_ins.columns[:-1]
importance1=pd.Series(rf.feature_importances_,index=index1)
importance1.nlargest(13).plot(kind='barh',color='pink')

elif max_==C1_:
    print('SELECTED ALGORITHM : NAIVE BAYES: ')
    print('TYPE OF ALGORITHM : REGRESSION')
    print('ACCURACY SCORE :',max_)
    print( )
    print('CONTRIBUTION')
    print( )

    importance2=BR.feature_importances_
    for i,v in enumerate(importance2):
        print("Feature: %0d ; Score: %5f"%(i,v))

    index2=med_ins.columns[:-1]
    importance2=pd.Series(rf.feature_importances_,index=index2)
    importance2.nlargest(13).plot(kind='barh',color='pink')

elif max_==D1_:

    print('SELECTED ALGORITHM : RANDOM FOREST ')
    print('TYPE OF ALGORITHM : REGRESSION')
    print('ACCUARACY SCORE :',max_)
    print( )
    print('CONTRIBUTION')
    print( )

    importance3=rf.feature_importances_
    for i,v in enumerate(importance3):
        print("Feature: %0d ; Score: %5f"%(i,v))

    index3=med_ins.columns[:-1]
    importance3=pd.Series(rf.feature_importances_,index=index3)
    importance3.nlargest(13).plot(kind='barh',color='pink')

elif max_==E1_:

    print('SELECTED ALGORITHM :KNN ')
```

```
print('TYPE OF ALGORITHM : REGRESSION')
print('ACCUARACY SCORE :',max_)
print( )
print('CONTRIBUTION')
print( )

importance4=kn.feature_importances_
for i,v in enumerate(importance4):
    print("Feature: %0d ; Score: %5f"%(i,v))

index4=med_ins.columns[:-1]
importance4=pd.Series(rf.feature_importances_,index=index4)
importance4.nlargest(13).plot(kind='barh',color='pink')

elif max_==F1_:
    print('SVM:',max_)
    print('SELECTED ALGORITHM : SVM ')
    print('TYPE OF ALGORITHM : REGRESSION')
    print('ACCUARACY SCORE :',max_)
    print( )
    print( )
    print('CONTRIBUTION')
    print( )

    importance5=sv.feature_importances_
    for i,v in enumerate(importance5):
        print("Feature: %0d ; Score: %5f"%(i,v))

    index5=med_ins.columns[:-1]
    importance5=pd.Series(rf.feature_importances_,index=index5)
    importance5.nlargest(13).plot(kind='barh',color='pink')

else:
    print('OOPS!! SOMETHING WENT WRONG')
```

ALGORITHM ACCURACY SCORES